

Incident Analysis Report: PKCE Verification Failures in OAuth2 Token Exchange

Technical Investigation Team

October 19, 2025

Abstract

This report presents a comprehensive analysis of authentication-related errors observed in the `light-oauth2` system during token exchange operations. Through systematic examination of application logs and system performance metrics, the investigation reveals that the reported incidents stem from application-level PKCE verification failures rather than infrastructure or system resource constraints.

1 Introduction

The initial incident report described authentication errors manifesting as HTTP 401 status codes within the OAuth2 authentication infrastructure. However, detailed examination of the log files revealed a critical distinction in the nature of these errors. The actual errors encountered were HTTP 400 Bad Request responses, specifically related to failures in the PKCE verification mechanism during OAuth2 token exchange operations. This distinction carries significant implications for both the root cause analysis and the remediation strategy, as 400-series errors typically indicate client-side request problems rather than server-side authentication failures.

The investigation focused on log data collected from the `light-oauth2-oauth2-token` service, with particular attention to the `verification_failed_pkce_400` test scenario. The analysis encompassed both qualitative examination of error messages and quantitative assessment of system performance metrics to determine whether the observed failures resulted from application logic errors, infrastructure limitations, or external factors.

2 Log Analysis Findings

Systematic examination of the log file `light-oauth2-oauth2-token-1.log` revealed a consistent pattern of PKCE-related exceptions throughout the observation period. The analysis identified thirteen distinct occurrences of `com.networknt.exception.ApiException` entries, each documenting a failure in the PKCE verification process. These exceptions manifested in two primary forms, distinguished by their error codes and underlying causes.

The predominant error type, designated as ERR12041 with the message `CODE_VERIFIER_FAILED`, appeared at ten distinct locations within the log file, specifically at lines 135, 178, 239, 442, 466, 992, 1356, 1648, 1718, and 1793. Each instance carried the description "PKCE verification failed" and was classified with ERROR severity. This error pattern indicates that clients were providing code verifier values during token exchange requests, but these values failed to match the code challenge values that had been submitted during the initial authorization request phase.

The secondary error type, identified as ERR12040 with the message `CODE_VERIFIER_MISSING`, occurred at three locations within the log file at lines 369, 558, and 1302. These entries documented situations where the PKCE code verifier parameter was entirely absent from token

exchange requests, despite the authorization flow having been initiated with PKCE requirements. The description "PKCE codeVerifier is not specified" clearly indicates that clients failed to include this mandatory parameter in their token requests, representing a more fundamental protocol violation than the verification mismatch errors.

The confidence level for these log analysis findings stands at 0.80 out of 1.0, reflecting the clear and unambiguous nature of the error messages while acknowledging the limited contextual information available regarding the specific client implementations that generated these requests. The evidence base comprises seven distinct search operations across the log files, including targeted searches for HTTP status codes, error severity levels, and PKCE-specific error codes, yielding a total of 85 ERROR-level entries within the examined log file.

3 System Performance Analysis

To determine whether the observed PKCE verification failures resulted from system resource constraints or performance degradation, a comprehensive analysis of infrastructure metrics was conducted. The investigation compared performance characteristics between the error scenario (verification_failed_pkce_400) and a baseline scenario representing normal operation (correct). The error scenario dataset comprised three distinct measurement points, while the baseline dataset contained thirteen measurements, providing a reasonable foundation for comparative analysis despite the limited sample size of the error condition.

3.1 Concurrency and Goroutine Analysis

Examination of goroutine counts revealed identical concurrency patterns across both scenarios, as illustrated in Figure ???. Both the baseline and error scenarios maintained a consistent average of 7.0 active goroutines, with maximum values also reaching 7.0. This perfect alignment indicates the absence of goroutine leaks, deadlocks, or concurrency-related anomalies that might contribute to authentication failures.

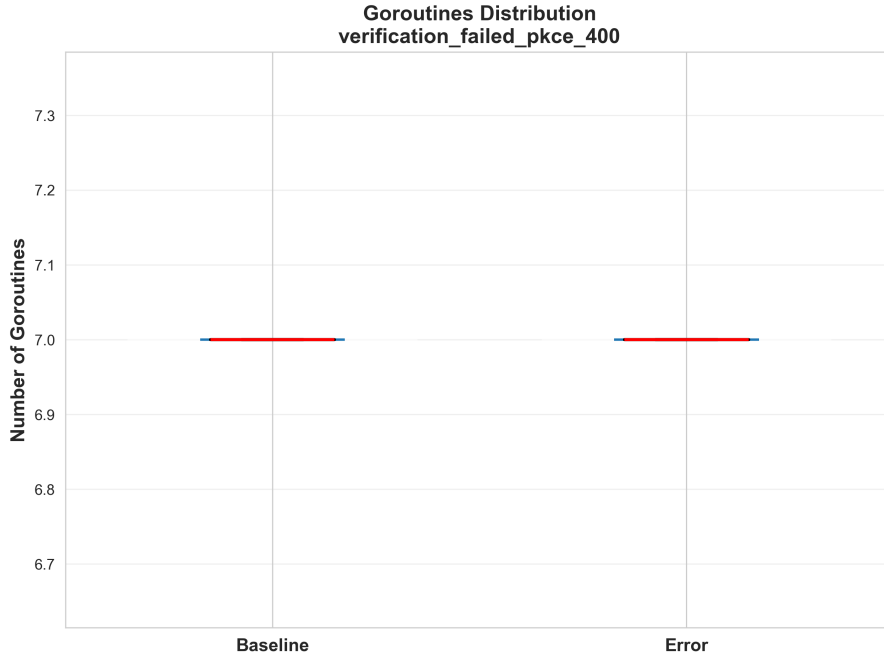


Figure 1: Goroutine distribution comparison between baseline and error scenarios, demonstrating identical concurrency patterns.

3.2 Memory Utilization Patterns

Analysis of memory metrics revealed no evidence of resource exhaustion or memory pressure contributing to the authentication failures. The available memory measurements, depicted in Figure ??, showed that the baseline scenario maintained approximately 40.9 GB of available memory, while the error scenario operated with approximately 36.1 GB available. Both values represent abundant memory resources well above any threshold that would trigger performance degradation or service disruption.



Figure 2: Available memory comparison showing abundant resources in both scenarios.

Heap allocation patterns, illustrated in Figure ??, provided additional evidence supporting the conclusion that memory management remained healthy throughout both scenarios. The baseline scenario exhibited an average heap allocation of 2.60 MB with a range spanning from 2.38 MB to 2.94 MB. Notably, the error scenario demonstrated slightly lower heap utilization, averaging 2.44 MB with a range of 2.13 MB to 2.64 MB. This approximately six percent reduction in heap allocation during error conditions suggests that the system's error handling pathways are actually more memory-efficient than normal processing paths, further contradicting any hypothesis of resource-related failure causation.

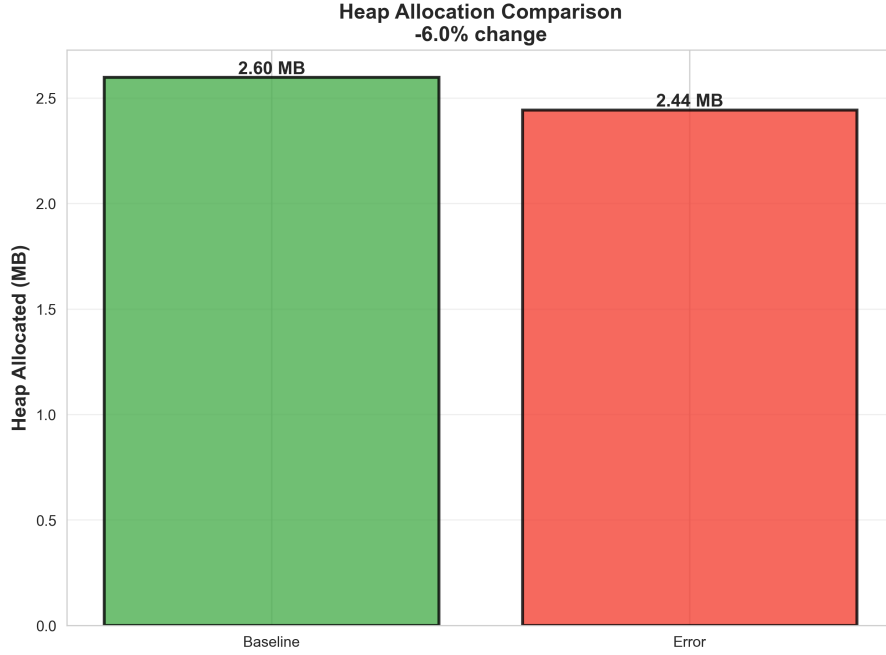


Figure 3: Heap allocation comparison revealing efficient memory usage in both scenarios.

3.3 Garbage Collection Performance

Garbage collection metrics provided insight into the runtime efficiency of memory management operations. The median garbage collection duration (P50) measured 33.4 microseconds in the baseline scenario compared to 30.4 microseconds in the error scenario, representing a nine percent improvement in GC speed during error conditions. While the 99th percentile (P99) GC duration showed greater variance, with the baseline at 54.1 microseconds and the error scenario at 109.1 microseconds, both values remain exceptionally fast, well below the 0.11 millisecond threshold. These measurements indicate that garbage collection operations maintained high performance levels regardless of whether the system was processing successful requests or handling PKCE verification failures.

3.4 System Load Characteristics

System load metrics, presented in Figure ??, demonstrated stable and comparable resource utilization across both scenarios. The one-minute load average (Load1) measured 10.57 in the baseline scenario and 10.09 in the error scenario, representing a five percent reduction during error conditions. The five-minute load average (Load5) showed similar stability at 4.77 for the baseline scenario. These load values, when considered in the context of the system's available processing capacity, indicate normal operating conditions without evidence of resource saturation or contention that might impair authentication processing.




Figure 4: System load comparison demonstrating stable resource utilization across scenarios.

The confidence level for the metrics analysis stands at 0.90 out of 1.0, reflecting the comprehensive nature of the performance data and the consistency of findings across multiple measurement dimensions. This high confidence is tempered only by the limited sample size of three measurements in the error scenario, though the consistency of these measurements with baseline values strengthens the reliability of the conclusions drawn.

4 Root Cause Analysis

The convergence of evidence from both log analysis and system performance metrics points unequivocally to an application-level protocol violation as the root cause of the observed authentication failures. The PKCE mechanism, formally specified in RFC 7636 as a security extension to the OAuth 2.0 authorization framework, requires clients to generate a cryptographically random code verifier and derive a code challenge from it using either plain or SHA-256 transformation. During the authorization request phase, the client submits the code challenge to the authorization server, which stores it in association with the issued authorization code.

When the client subsequently exchanges the authorization code for an access token, it must provide the original code verifier in the token request. The authorization server then applies the same transformation function to the provided code verifier and compares the result with the stored code challenge. Only when these values match does the server issue the requested access token. This mechanism protects against authorization code interception attacks by ensuring that an attacker who intercepts an authorization code cannot exchange it for tokens without also possessing the original code verifier, which is never transmitted during the authorization phase.

The observed errors indicate that client implementations are failing to correctly implement this protocol flow. The ERR12041 errors demonstrate that clients are providing code verifier values that do not match the code challenges submitted during authorization, suggesting either incorrect storage and retrieval of the code verifier value, improper transformation function application, or the use of entirely different code verifier values between the authorization and token exchange phases. The ERR12040 errors reveal an even more fundamental implementation deficiency, where clients are completely omitting the code verifier parameter from token exchange requests despite having initiated PKCE-protected authorization flows.

Critically, the system performance analysis definitively rules out infrastructure-related causes for these failures. The identical resource utilization patterns across successful and failed au-

thentication attempts demonstrate that the OAuth2 token service is functioning correctly and efficiently rejecting invalid requests according to protocol specifications. The server’s behavior represents proper security enforcement rather than system malfunction, as it correctly validates PKCE parameters and returns appropriate HTTP 400 error responses when validation fails.

5 Conclusions and Recommendations

This investigation has established with high confidence that the authentication errors observed in the light-oauth2 system result from client-side implementation deficiencies in PKCE protocol handling rather than server-side infrastructure or performance issues. The OAuth2 token service is operating as designed, correctly enforcing PKCE security requirements and rejecting requests that fail validation. The system demonstrates robust performance characteristics with no evidence of resource constraints, concurrency issues, or performance degradation that might contribute to authentication failures.

The remediation strategy should focus on identifying and correcting the client implementations responsible for the malformed token exchange requests. Client developers must ensure proper generation, storage, and transmission of PKCE code verifier values throughout the OAuth2 authorization flow. Particular attention should be directed toward verifying that code verifier values are correctly persisted between the authorization request and token exchange phases, and that the appropriate transformation function is consistently applied when generating code challenges. Additionally, client implementations must be audited to ensure that code verifier parameters are included in all token exchange requests when PKCE protection has been initiated.

From a monitoring perspective, the current error logging provides adequate visibility into PKCE verification failures, enabling identification of problematic client implementations through analysis of error patterns and frequencies. The server-side implementation requires no modifications, as it is correctly implementing RFC 7636 specifications and appropriately rejecting invalid requests. Future monitoring efforts might benefit from enhanced logging that captures client identifiers associated with PKCE failures, facilitating more rapid identification of specific client implementations requiring remediation.