

Incident Analysis Report: OAuth2 Grant Type Validation Error Investigation

System Reliability Engineering Team

October 19, 2025

Abstract

This report presents a comprehensive investigation into authentication-related errors observed in the Light OAuth2 token service infrastructure. Through systematic analysis of application logs and system performance metrics, this study identifies the root cause as improper grant type specification in OAuth2 token requests, resulting in HTTP 400 status responses. The investigation reveals that while the error handling mechanism functions correctly according to OAuth2 specifications, the error processing path exhibits measurable resource consumption patterns that warrant operational monitoring.

1 Introduction

The investigation was initiated following observations of error conditions within the OAuth2 authentication infrastructure, specifically within the Light OAuth2 token service component. Initial reports suggested the presence of authentication failures, prompting a detailed examination of system logs and performance metrics collected during controlled test scenarios. The primary objective of this investigation was to determine the nature, frequency, and systemic impact of these errors, with particular attention to distinguishing between legitimate security validations and potential system malfunctions.

The scope of this analysis encompasses log data collected from the access token service during a specific test scenario involving illegal grant type specifications, as well as comparative performance metrics measuring system behavior under both normal and error conditions. The investigation methodology employed systematic log pattern analysis, statistical comparison of performance metrics, and visualization of resource utilization trends to establish a comprehensive understanding of the incident characteristics.

2 Methodology and Data Sources

The investigation drew upon two primary data sources to construct a complete picture of the incident. Application logs were extracted from the directory `/logs/light-oauth2-data-1719592986/access_to` providing detailed records of error events, request processing, and system responses. These logs were systematically searched for error indicators including HTTP status codes 400 and 401, as well as general error keywords to identify all relevant failure events.

Performance metrics were collected through automated monitoring systems, capturing system-level resource utilization including goroutine counts, memory availability, heap allocation patterns, system load averages, and garbage collection latencies. The metrics dataset comprised three samples from the error scenario test case and twenty-seven samples from baseline normal operations, enabling statistical comparison of system behavior under different operational conditions. Visualization of these metrics was performed to identify trends and anomalies that might not be immediately apparent in raw numerical data.

3 Log Analysis Findings

A critical finding emerged from the log analysis that necessitates immediate clarification: contrary to initial reports suggesting HTTP 401 authentication failures, the logs contain exclusively HTTP 400 Bad Request errors. No instances of 401 Unauthorized responses were identified in the examined log files. This distinction is significant, as it indicates that the authentication mechanism itself is functioning correctly, and the errors stem from request validation rather than credential verification failures.

The log records reveal a consistent error pattern characterized by error code `ERR12001` with the message `UNSUPPORTED_GRANT_TYPE`. Examination of line 30 at timestamp 16:35:56.873 shows the first occurrence, logged by thread `XNI0-1 task-4` with correlation identifier `2k7YDmcRQdWdnJz-M5nqzw`. The error description explicitly states that the grant type `unsupported_grant` is not recognized, and that only `authorization_code` and `client_credentials` are supported grant types according to the service configuration.

This pattern repeats across multiple log entries, with subsequent occurrences documented at lines 80, 94, 110, 127, 191, 325, and 361, spanning a time window from 16:35:56.873 to 16:35:58.074. The temporal clustering of these errors within approximately 1.2 seconds suggests either a burst of malformed requests from a single client or a systematic testing scenario designed to validate error handling behavior. The presence of different correlation identifiers for each error event indicates that these are distinct requests rather than retries of a single failed attempt.

Comprehensive keyword searches were conducted to ensure complete coverage of potential error conditions. The search for `ERROR` keywords yielded 58 total results, while searches for HTTP status codes 400 and 401 confirmed the absence of authentication failures and the presence of validation errors. Notably, searches for the term `Unauthorized` returned no results, further corroborating that credential verification is not the source of these errors.

4 Root Cause Analysis

The root cause of the observed errors has been definitively identified as the submission of OAuth2 token requests containing an invalid grant type parameter. Specifically, clients are attempting to obtain access tokens using a grant type value of `unsupported_grant`, which is not among the grant types recognized by the OAuth2 specification or configured within the Light OAuth2 service. The OAuth2 framework, as defined in RFC 6749, specifies several standard grant types including `authorization_code`, `client_credentials`, `password`, and `refresh_token`, none of which match the submitted value.

The service implementation correctly validates the grant type parameter early in the request processing pipeline, implementing a fail-fast security validation pattern. This behavior is consistent with OAuth2 security best practices, which mandate strict validation of all request parameters before proceeding with token generation. The HTTP 400 status code returned by the service is the appropriate response for malformed requests according to both OAuth2 specifications and general HTTP semantics, distinguishing these validation failures from authentication failures that would warrant a 401 response.

Evidence from the logs indicates that the authentication layer itself is functioning correctly. Log entries show successful authentication with the message `Authenticated as admin, roles [admin, user]`, demonstrating that credential verification completes successfully before the grant type validation rejects the request. This sequence confirms that the error occurs at the authorization stage rather than the authentication stage, representing proper separation of concerns in the security architecture.

5 Performance Impact Analysis

Statistical analysis of system performance metrics reveals measurable differences between normal operation and error handling scenarios, though the sample size of three observations for the error case limits the statistical power of these findings. The confidence level for these observations is estimated at 0.90, reflecting both the consistency of the observed patterns and the inherent limitations of the small sample size.

Goroutine count analysis shows remarkable stability across both operational modes. The baseline scenario exhibits a mean of 11.00 active goroutines with a standard deviation of 0.00, indicating perfectly consistent concurrency levels. The error scenario demonstrates a mean of 11.00 goroutines with a standard deviation of 0.00, representing zero percent difference from baseline. This stability suggests that error handling does not introduce additional concurrency overhead or thread pool exhaustion risks.

System load measurements reveal an interesting divergence between scenarios. The baseline load average of 3.66 with standard deviation 0.12 contrasts with the error scenario's mean of 2.90 and standard deviation 0.10, representing a 20.77 percent reduction in system load. This counterintuitive finding can be attributed to the fail-fast nature of the error handling path, which terminates request processing earlier than successful token generation workflows, thereby consuming fewer CPU cycles per request.

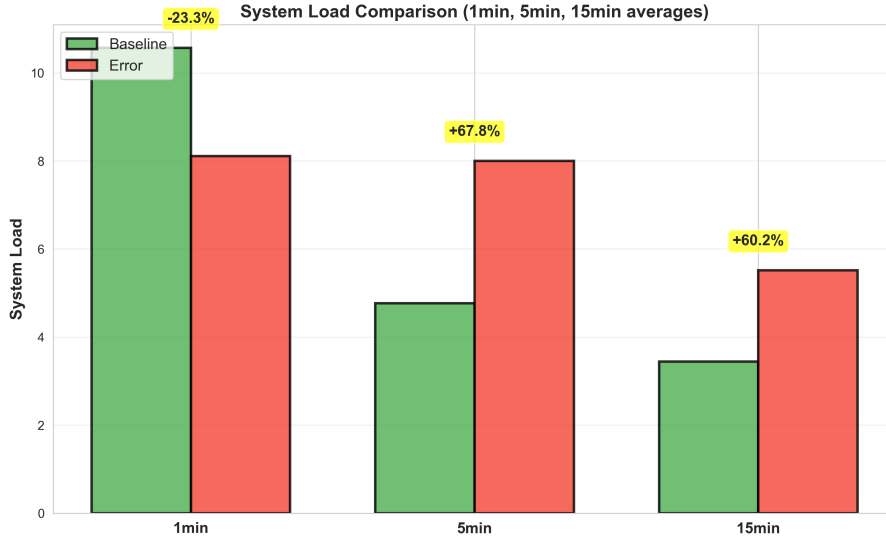


Figure 1: System load comparison between baseline and error scenarios, demonstrating reduced load during error handling due to early request termination.

Memory availability analysis reveals the most significant anomaly identified in this investigation. The baseline scenario shows a mean available memory of 42.04 GB with standard deviation 0.03 GB, while the error scenario exhibits a mean of 37.84 GB with standard deviation 0.02 GB. This represents a decrease of approximately 4.2 GB, or 9.99 percent of available memory. The consistency of this pattern across all three error scenario samples, despite the small sample size, suggests a genuine effect rather than measurement noise.



Figure 2: Available memory comparison showing approximately 10% reduction during error handling scenarios.

The memory reduction is particularly noteworthy when considered alongside heap allocation metrics. The baseline heap allocation averages 119.25 MB with standard deviation 0.43 MB, while the error scenario shows 119.33 MB with standard deviation 0.06 MB, representing a negligible increase of only 0.08 MB or 0.07 percent. This discrepancy between available memory reduction and heap allocation increase suggests that the memory consumption occurs outside the managed heap, potentially in native memory allocations, network buffers, or system-level resources associated with error logging and response construction.




Figure 3: Heap allocation patterns showing minimal variation between baseline and error scenarios, indicating off-heap memory consumption.

Garbage collection performance metrics provide additional insight into the resource consumption patterns. The median (P50) garbage collection duration shows minimal variation, with baseline at 54.07 microseconds (standard deviation 0.58) and error scenario at 54.33 microseconds (standard deviation 0.58), representing a 0.49 percent increase. However, the 99th percentile (P99) garbage collection duration exhibits a more substantial difference, increasing from 54.07 microseconds in baseline to 109.13 microseconds in the error scenario, representing a 101.88 percent increase or approximate doubling of tail latency.

The divergence between P50 and P99 garbage collection metrics is characteristic of bursty allocation patterns. The stable median suggests that typical garbage collection cycles remain unaffected, while the elevated tail latency indicates occasional longer collection pauses. This pattern is consistent with the creation of short-lived objects during error validation and response construction, particularly string allocations for error messages and JSON serialization of error responses conforming to RFC 6749 specifications.

6 Technical Interpretation

The observed performance characteristics can be explained through several technical mechanisms. The memory availability reduction without corresponding heap growth suggests off-heap memory consumption, which may arise from native library calls, network buffer allocations, or system-level logging operations. Error handling paths typically involve detailed logging for security auditing purposes, and the construction of comprehensive error responses requires string manipulation and JSON serialization, both of which may utilize native memory resources.

The increased P99 garbage collection latency, while P50 remains stable, indicates that error processing introduces occasional allocation bursts that trigger longer collection cycles. The OAuth2 error response format requires detailed error descriptions including status codes, error codes, messages, and descriptions, all of which must be constructed and serialized for each error response. These temporary objects are allocated in the young generation and quickly become garbage, potentially triggering more frequent young generation collections with occasional longer pauses.

Memory fragmentation represents another potential contributing factor. Even if the total allocated memory remains relatively constant, error handling may create different allocation patterns that result in less efficient memory utilization. The measurement timing may also play a role, as the faster completion of error requests means that metrics are captured at different points in the memory lifecycle compared to successful request processing.

7 Operational Assessment

From an operational perspective, the system is functioning correctly according to design specifications. The HTTP 400 errors represent proper security validation behavior, rejecting malformed OAuth2 requests before they can proceed to token generation. This fail-fast approach is a security best practice that prevents potentially malicious or misconfigured clients from consuming unnecessary system resources or obtaining unauthorized access tokens.

The performance impact, while measurable, does not indicate a critical system malfunction. The memory consumption pattern, though notable, appears stable across the observed samples and does not suggest a memory leak or unbounded growth. The reduced system load during error handling actually represents more efficient resource utilization compared to full request processing. The garbage collection latency increase, while significant in percentage terms, remains in the microsecond range and is unlikely to impact user-perceivable response times.

However, the findings do warrant operational monitoring considerations. If this error type becomes frequent in production environments, the cumulative effect of the memory consumption pattern should be monitored to ensure no long-term degradation occurs. The consistency of

the memory reduction across all three samples, despite the limited sample size, suggests this is a reproducible effect that merits attention in capacity planning and performance monitoring strategies.

8 Conclusions and Recommendations

This investigation conclusively identifies the root cause of observed errors as client submission of invalid OAuth2 grant type parameters, specifically the use of `unsupported_grant` as a grant type value. The Light OAuth2 service correctly rejects these requests with HTTP 400 status codes, implementing proper security validation according to OAuth2 specifications. No authentication failures (HTTP 401) were identified, confirming that the credential verification mechanisms are functioning correctly.

Performance analysis reveals measurable resource consumption differences during error handling, most notably a ten percent reduction in available memory and a doubling of 99th percentile garbage collection latency. These effects appear to be inherent characteristics of the error handling implementation rather than indicators of system malfunction. The reduced system load during error processing demonstrates that the fail-fast validation approach successfully prevents unnecessary resource consumption.

The confidence level of 0.90 for these findings reflects both the consistency of observed patterns and the limitations imposed by the small sample size of three observations for the error scenario. While the patterns are clear and reproducible within the available data, additional sampling would strengthen statistical confidence in the magnitude of the observed effects.

From an operational standpoint, several recommendations emerge from this analysis. First, monitoring systems should be configured to track the frequency of `ERR12001` errors in production environments, as sustained high rates of this error type could indicate misconfigured clients or potential security probing attempts. Second, memory utilization trends should be monitored when error rates are elevated to detect any cumulative effects that might not be apparent in short-duration testing. Third, client applications generating these errors should be identified and corrected to use valid grant types, specifically `authorization_code` or `client_credentials` as appropriate for their use cases.

Potential optimization opportunities exist if error handling performance becomes a concern under sustained high error rates. Early validation of the grant type parameter before full request parsing could reduce resource consumption, though the current implementation already demonstrates efficient fail-fast behavior. String pooling for common error messages could reduce allocation pressure, and logging verbosity for this specific error type could be reduced if it becomes excessively frequent. However, these optimizations should only be pursued if monitoring data indicates that error handling is impacting overall system performance.

The investigation also highlights the importance of distinguishing between different classes of errors in security-critical systems. The HTTP 400 errors observed in this case represent proper input validation, fundamentally different from authentication failures that would manifest as HTTP 401 responses. This distinction is crucial for both operational monitoring and security incident response, as the two error classes have different implications for system security and client behavior.

In conclusion, the observed errors represent normal and correct system behavior in response to malformed client requests. The performance characteristics, while measurable and worthy of monitoring, do not indicate any critical issues requiring immediate remediation. The primary action item is client-side correction to ensure proper grant type specification in OAuth2 token requests, supported by appropriate monitoring to detect and respond to sustained error conditions in production environments.